


Class 07: Machine Learning 1

Qingyun Zheng (A16338254)

Today we will begin our exploration of some “classical” machine learning approaches. We’ll start with clustering:

Let’s first make up some data cluster where we know what the answer should be


```
hist(rnorm(1000))
```



```
x <- c(rnorm(30, -3), rnorm(30, 3))
y <- rev(x)
z <- cbind(x,y)
```

Let’s plot z using `plot()`

```
plot(z)
```



The main function in “base” R for k-means clustering is called `kmeans()`

```
k <- kmeans(z, centers = 2)  
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```

          x      y
1 -2.782877 3.113843
2 3.113843 -2.782877

```

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 60.3723 60.3723
```

(between_SS / total_SS = 89.6 %)

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. How big are the clusters (i.e. their size)?

k\$size


[1] 30 30

Q. What clusters do my data points reside in?

k\$cluster


Q. Make a plot of our data of colored by cluster assignment - make a result figure

```
plot(z, col = k$cluster)
points(k$centers, col = "blue", pch = 15)
```



Q. Clustering with k-means into 4 clusters and plot your results as above

```
k4 <- kmeans(z, centers = 4)
plot(z, col = k4$cluster)
points(k4$centers, col = "blue", pch = 15)
```




Run K-means with centers (values of k) equal to 1 to 6

```
ans <- NULL
for(i in 1:6){
  ans <- c(ans, kmeans(z, centers = i)$tot.withinss)
}
ans
```

```
[1] 1163.88386 120.74459 96.24586 71.74712 59.64421 75.72588
```

```
plot(ans, typ = "b")
```



Hierachial clustering

The main function in “base” R for this is called `hclust()`


```
d <- dist(z)
hc <- hclust(d)
hc
```

Call:
`hclust(d = d)`

Cluster method : complete
Distance : euclidean
Number of objects: 60

```
plot(hc)
abline(h=7, col = "red")
```


Cluster Dendrogram



```
d  
hclust (*, "complete")
```

To obtain clusters from hour `hclust` result object `hc` we “cut the tree to yield different sub branches. For this we use the `cutree()` function

```
grps <- cutree(hc, h=7)  
plot(z, col= grps)
```



Principal component Analysis (PCA)

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17 5
```

There are 17 rows and 5 columns

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139


The second method with `read.csv()` is more robust. When you run the first method, i.e. indexing, the index changes everytime and we lose one column until it run out of the dimension of the df.

Spotting major differences and trends

```
rainbow(nrow(x))
```


```
[1] "#FF0000" "#FF5A00" "#FFB400" "#FOFF00" "#96FF00" "#3CFF00" "#00FF1E"
[8] "#00FF78" "#00FFD2" "#00D2FF" "#0078FF" "#001EFF" "#3C00FF" "#9600FF"
[15] "#F000FF" "#FF00B4" "#FF005A"
```

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above barplot() function results in the following plot?

```
barplot(as.matrix(x), beside=FALSE, col=rainbow(nrow(x)))
```




By setting the `beside = FALSE` we can make the stacked barplot

Q4: Changing what optional argument in the above `ggplot()` code results in a stacked barplot figure?

Paris plots and heatmaps

Scatterplot matrices (aka “paris plots”) can be useful for the relatively small datasets like this one.


```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



Q5: We can use the `pairs()` function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

The x and y axes are different countries. For example, England on x axis and Wales on y axis. The points represent different food. So the plots are showing the amount of food consumed in one country vs in another country. If a point lies on the diagonal then it means the quantity of this food consumed in both countries is the same.

```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

It looks like Wales and England are quite similar in their consumption of these foods, It is still quite difficult to tell what is going on in the dataset

PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`

As we want to do PCA on the food data for the different countries we will want the foods in the columns

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

our result object is called `pca` and it has a `$x` component that we will look at first


```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points. Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
library(ggplot2)
cols <- c("orange", "red", "blue", "darkgreen")

ggplot(pca$x) +
  aes(PC1, PC2, label = rownames(pca$x)) +
  geom_point(col = cols) +
  geom_text(vjust = -0.5) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()
```



Another major result out of PCA is the “variable loadings” or `$rotation` that tells us how the original variables (foods) contribute to the PCs (our new axes)

```
og <- pca$rotation
ggplot(og) +
  aes(PC1, rownames(og)) +
  geom_col()
```

